


Rates of convergence in Bayesian meta-learning

Pierre Alquier



12th Workshop on High Dimensional Data Analysis
ESSEC and University Mohamed V
Rabat, Morocco



Charles RIOU

University of Tokyo,
RIKEN AIP



Badr-Eddine CHÉRIEF-ABDELLATIF


CNRS,
Sorbonne Université



Riou, C., Alquier, P. and Chérif-Abdellatif, B.-E. (2023). Bayes meets Bernstein at the Meta Level : an Analysis of Fast Rates in Meta-Learning with PAC-Bayes. Preprint arXiv :2302.11709.


- 1 Introduction : Bayesian learning and meta-learning
- 2 Overview of our results
- 3 More detailed view of our results

- 1 Introduction : Bayesian learning and meta-learning
- 2 Overview of our results
- 3 More detailed view of our results



$$\pi(\theta|x_1, \dots, x_n)$$

$$\propto \pi(\theta) \prod_{i=1}^n p(x_i|\theta)$$



$$\rho(\theta)$$

$$\propto \pi(\theta) e^{-\alpha \sum_{i=1}^n \ell(x_i, \theta)}$$


To keep the results as simple as possible :

- $\mathcal{S} = (X_1, \dots, X_n)$ i.i.d. from P ,
- $\ell(x, \theta)$ bounded by 1.
- Generalization risk : $R(\theta) = \mathbb{E}_{X \sim P}[\ell(X, \theta)]$.
- Objective : $\theta^* = \arg \min_{\theta \in \Theta} R(\theta)$.
- Risk of the “Bayes” procedure $\rho : \mathbb{E}_{\theta \sim \rho}[R(\theta)]$.

Theorem (stated informally)

$$\mathbb{E}_{\mathcal{S}} \left\{ \mathbb{E}_{\theta \sim \rho}[R(\theta)] \right\} \leq R(\theta^*) + c \sqrt{\frac{d \log(n)}{n}}$$

where $d = d(P, \pi)$ defined in the next slide, for α well chosen.



$$N(\theta^*, s) := \{\theta \in \Theta : R(\theta) - R(\theta^*) \leq s\}.$$

$d = d(P, \pi)$ is the smallest number such that, for any s small enough :

$$\pi(N(\theta^*, s)) \geq s^d.$$

How do we prove the theorem ?

$$\begin{aligned}\rho(\theta) &\propto \pi(\theta) e^{-\alpha \sum_{i=1}^n \ell(x_i, \theta)} \\ &= \arg \min_{p \in \mathcal{P}(\Theta)} \left\{ \mathbb{E}_{\theta \sim p} \left[\frac{1}{n} \sum_{i=1}^n \ell(x_i, \theta) \right] + \frac{KL(p, \pi)}{\alpha n} \right\}.\end{aligned}$$

PAC-Bayes / Information bounds

$$\mathbb{E}_{\mathcal{S}} \left\{ \mathbb{E}_{\theta \sim \rho} [R(\theta)] \right\} \leq \inf_p \left\{ \mathbb{E}_{\theta \sim p} [R(\theta)] + \alpha + \frac{KL(p, \pi)}{\alpha n} \right\}.$$

In particular, for p as the restriction of π to $N(\theta^*, s)$,

$$\mathbb{E}_{\mathcal{S}} \left\{ \mathbb{E}_{\theta \sim \rho} [R(\theta)] \right\} \leq \inf_{s > 0} \left\{ R(\theta^*) + s + \alpha + \frac{d \log \frac{1}{s}}{\alpha n} \right\}.$$

- Old result : in a “noiseless setting”, when there is a θ such that $\ell(x, \theta) = 0$ almost surely for $x \sim P$,

$$\mathbb{E}_{\mathcal{S}} \left\{ \mathbb{E}_{\theta \sim \rho} [R(\theta)] \right\} \leq \underbrace{R(\theta^*)}_{=0} + c \frac{d \log(n)}{n}.$$

- Similar fast rates obtained in classification under Mammen and Tsybakov margin assumption (1999).
- Also with Lipschitz and strongly convex losses $\ell(x, \cdot)$ by Bartlett and Mendelson (2006).

All these assumptions turned out to be a special case of :

Bernstein condition

$$\mathbb{E}_{x \sim P} \left\{ [\ell(x, \theta) - \ell(x, \theta^*)]^2 \right\} \leq C [R(\theta) - R(\theta^*)].$$

What about variational Bayes?

Let \mathcal{W} be a subset of $\mathcal{P}rob(\Theta)$, and put :

$$\rho^{\mathcal{W}}(\theta) = \arg \min_{p \in \mathcal{P}rob(\Theta) \cap \mathcal{W}} \left\{ \mathbb{E}_{\theta \sim p} \left[\frac{1}{n} \sum_{i=1}^n \ell(x_i, \theta) \right] + \frac{KL(p, \pi)}{\alpha n} \right\}.$$



P. Alquier, J. Ridgway, N. Chopin (2016). On the Properties of Variational Approximations of Gibbs Posteriors. JMLR.

provides minimal assumptions on \mathcal{W} ensuring

$$\mathbb{E}_{\mathcal{S}} \left\{ \mathbb{E}_{\theta \sim \rho^{\mathcal{W}}} [R(\theta)] \right\} \leq R(\theta^*) + c \left(\frac{d(P, \pi) \log(n)}{n} \right)^{\beta}$$

where $\beta = 1$ under Bernstein condition, $\beta = 1/2$ otherwise.

Recap

$$\rho(\theta) \propto \pi(\theta) e^{-\alpha \sum_{i=1}^n \ell(x_i, \theta)}.$$


We have :

$$\mathbb{E}_{\mathcal{S}} \left\{ \mathbb{E}_{\theta \sim \rho} [R(\theta)] \right\} \leq R(\theta^*) + c \left(\frac{d \log(n)}{n} \right)^{\beta}$$

where $\beta = 1$ under Bernstein condition, $\beta = 1/2$ otherwise.

- The generalization error is driven by $d = d(P, \pi)$ that depends on π .
- Tempting to learn a better π , but π is not allowed to depend on the data...

Idea of Bayesian meta-learning :



- We solve many related tasks (say T) using Bayesian learning.
- By related, we mean that the same prior could be used in all tasks.
- Based on past tasks, can we define a π that would work better for future tasks?


Notations :


- Tasks : $t = 1, \dots, T$.
- P_1, \dots, P_T are i.i.d from \mathcal{P} .
- Task t : $\mathcal{S}_t = (x_{t,1}, \dots, x_{t,n})$ i.i.d from P_t .
- Generalization error in task t : $R_t(\theta) = \mathbb{E}_{x \sim P_t}[\ell(x, \theta)]$.
- Best error in task t : $R_t(\theta_t^*) = \min_{\theta} R_t(\theta)$.
- $\rho_t(\pi, \alpha)(\theta) \propto \pi(\theta) \exp[-\alpha \sum_{i=1}^n \ell(\theta, x_{t,i})]$.

Objective

- Learn $\hat{\pi} = \hat{\pi}(\mathcal{S}_1, \dots, \mathcal{S}_T)$.
- For a new task $P_{T+1} \sim \mathcal{P}$, $\mathcal{S}_{T+1} = (x_{T+1,1}, \dots, x_{T+1,n})$ i.i.d. from P_{T+1} , we want :

$$\mathbb{E}_{\theta \sim \rho_{T+1}(\hat{\pi}, \alpha)} [R_{T+1}(\theta)] \leq \mathbb{E}_{\theta \sim \rho_{T+1}(\pi, \alpha)} [R_{T+1}(\theta)] .$$





- Past tasks, used to learn a better prior. Expectation with respect to $P_1, \dots, P_T, \mathcal{S}_1, \dots, \mathcal{S}_T$ denoted by \mathbb{E}_{data} .
- New task. Expectation with respect to P_{T+1} and \mathcal{S}_{T+1} will be denoted by \mathbb{E}_{new} .

- 1 Introduction : Bayesian learning and meta-learning
- 2 Overview of our results
- 3 More detailed view of our results

Ultimate (non-achievable) performance :

$$\mathcal{E}^* = \mathbb{E}_{\text{new}}[R_{T+1}(\theta_{T+1}^*)] = \mathbb{E}_{P_{T+1} \sim \mathcal{P}}[R_{T+1}(\theta_{T+1}^*)].$$

With a fixed prior :

$$\begin{aligned}\mathcal{E}(\pi) &= \mathbb{E}_{\text{new}} \left\{ \mathbb{E}_{\theta \sim \rho_{T+1}(\pi, \alpha)} [R_{T+1}(\theta)] \right\} \\ &\leq \mathcal{E}^* + c \mathbb{E}_{P_{T+1} \sim \mathcal{P}} \left[\left(\frac{d(P_{T+1}, \pi) \log(n)}{n} \right)^\beta \right].\end{aligned}$$

To give an overview of our results, let us consider first an easy situation : we want to find the best of K priors, say

$$\pi_1, \dots, \pi_K.$$

Recall :

$$\rho_t(\pi, \alpha) = \arg \min_p \underbrace{\left\{ \mathbb{E}_{\theta \sim p} \left[\frac{1}{n} \sum_{i=1}^n \ell(x_{t,i}, \theta) \right] + \frac{KL(p, \pi)}{\alpha n} \right\}}_{\hat{\mathcal{R}}_t(p, \pi)}.$$

In this case, our procedure boils down to :

$$\hat{\pi} = \arg \min_{\pi \in \{\pi_1, \dots, \pi_K\}} \left\{ \frac{1}{T} \sum_{t=1}^T \hat{\mathcal{R}}_t \left[\rho_t(\pi, \alpha), \pi \right] \right\}.$$

Theorem

$$\begin{aligned}
& \mathbb{E}_{\text{data}}[\mathcal{E}(\hat{\pi})] \\
& \leq \min_{k=1,\dots,K} \mathcal{E}(\pi_k) + \frac{\log K}{T} \\
& \leq \mathcal{E}^* + c \min_{k=1,\dots,K} \mathbb{E}_{P_{T+1} \sim \mathcal{P}} \left[\left(\frac{d(P_{T+1}, \pi_k) \log(n)}{n} \right)^\beta \right] + \frac{\log K}{T}.
\end{aligned}$$

Important observations :

- gain expected only if $T \gg n$.
- the rate for learning the prior is in $1/T$ regardless of the rate within tasks ($\beta = 1$ or $\beta = 1/2$).

- More generally, we can learn the best prior in an infinite set \mathcal{Q} (for example, all Gaussian priors, etc).
- The definition of $\hat{\pi}$ gets a little more convoluted.
- We will recover similar results

$$\mathbb{E}_{\text{data}}[\mathcal{E}(\hat{\pi})] \leq \min_{\pi \in \mathcal{Q}} \mathcal{E}(\pi) + \frac{\mathcal{C}(\mathcal{Q})}{T}$$

where $\mathcal{C}(\mathcal{Q})$ is a complexity measure of \mathcal{Q} .

Example 1 : Gaussian priors.

- $\theta \in \mathbb{R}^p$.
- $\mathcal{Q} = \{\mathcal{N}(\mu, \Sigma), \mu \in \mathbb{R}^p, \Sigma \in \mathcal{S}_+^p\}$.
- fix some m and put $V = \mathbb{E}_{\text{new}} [\|\theta_{T+1}^* - m\|^2]$.

Very approximately,

$$\mathbb{E}_{\text{data}}[\mathcal{E}(\hat{\pi})] \leq \mathcal{E}^* + \frac{p}{T} + \begin{cases} \frac{V+p \log n}{n} & \text{if } V > \frac{n}{T} \\ 0 & \text{otherwise .} \end{cases}$$

Example 2 : mixture of Gaussian priors.

- $\theta \in \mathbb{R}^p$.
- $\mathcal{Q} = \left\{ \sum_{k=1}^K p_k \mathcal{N}(\mu_k, \Sigma_k) \right\}$.
- fix m_1, \dots, m_K and put $V = \mathbb{E}_{\text{new}} \left[\min_k \|\theta_{T+1}^* - m_k\|^2 \right]$.

$$\mathbb{E}_{\text{data}}[\mathcal{E}(\hat{\pi})] \leq \mathcal{E}^* + \frac{pK}{T} + \frac{\log K}{n} + \begin{cases} \frac{V + p \log n}{n} & \text{if } V > \frac{n}{T} \\ 0 & \text{otherwise .} \end{cases}$$

- 1 Introduction : Bayesian learning and meta-learning
- 2 Overview of our results
- 3 More detailed view of our results

The general procedure $\hat{\pi} = \hat{\pi}(\mathcal{S}_1, \dots, \mathcal{S}_T)$ is a little more convoluted, it is actually a Bayesian procedure :

- fix a prior Π on the set of priors $\mathcal{Q} : \Pi \in \mathcal{P}rob(\mathcal{Q})$,
- define :

$$\hat{\Lambda} = \arg \min_{\Lambda \in \mathcal{P}rob(\mathcal{Q})} \left\{ \mathbb{E}_{\pi \sim \Lambda} \left[\frac{1}{T} \sum_{t=1}^T \hat{\mathcal{R}}_t \left(\rho_t(\pi, \alpha), \pi \right) \right] + \frac{KL(\Lambda, \Pi)}{\gamma T} \right\},$$

- draw $\hat{\pi} \sim \hat{\Lambda}$.

Define

$$\pi^* = \arg \min_{\pi} \mathbb{E}_{\text{new}} \left[\hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi, \alpha), \pi \right) \right].$$

Lemma – Bernstein condition at the meta-level

For any $\pi \in \mathcal{Q}$,

$$\begin{aligned} & \mathbb{E}_{\text{new}} \left[\left(\hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi, \alpha), \pi \right) - \hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi^*, \alpha), \pi^* \right) \right)^2 \right] \\ & \leq C \mathbb{E}_{\text{new}} \left[\hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi, \alpha), \pi \right) - \hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi^*, \alpha), \pi^* \right) \right]. \end{aligned}$$

Theorem

$$\begin{aligned} \mathbb{E}_{\text{data}} \left\{ \mathbb{E}_{\hat{\pi} \sim \hat{\Lambda}} [\mathcal{E}(\hat{\pi})] \right\} &\leq \mathcal{E}^* \\ &+ \min_{\Lambda \in \text{Prob}(\mathcal{Q})} \mathbb{E}_{\pi \sim \Lambda} \left\{ \mathbb{E}_{P_{T+1} \sim \mathcal{P}} \left[\left(\frac{d(P_{T+1}, \pi) \log(n)}{n} \right)^\beta \right] \right. \\ &\quad \left. + \frac{\mathcal{K}(\Lambda, \Pi)}{\gamma T} \right\}. \end{aligned}$$

The aforementioned examples are obtained by specification of Π , and taking an explicit Λ above.

Remark :

$$\hat{\Lambda} = \arg \min_{\Lambda \in \mathcal{Prob}(\mathcal{Q})} \left\{ \mathbb{E}_{\pi \sim \Lambda} \left[\frac{1}{T} \sum_{t=1}^T \hat{\mathcal{R}}_t \left(\rho_t(\pi, \alpha), \pi \right) \right] + \frac{KL(\Lambda, \Pi)}{\gamma T} \right\}.$$

What happens if we minimize over a smaller set
 $\mathcal{V} \subset \mathcal{Prob}(\mathcal{Q})$?

$$\hat{\Lambda}_{\mathcal{V}} = \arg \min_{\Lambda \in \mathcal{V}} \left\{ \mathbb{E}_{\pi \sim \Lambda} \left[\frac{1}{T} \sum_{t=1}^T \hat{\mathcal{R}}_t \left(\rho_t(\pi, \alpha), \pi \right) \right] + \frac{KL(\Lambda, \Pi)}{\gamma T} \right\}.$$

Note : can be seen as a variational Bayes version of $\hat{\Lambda}$.

Theorem

$$\begin{aligned} \mathbb{E}_{\text{data}} \left\{ \mathbb{E}_{\hat{\pi} \sim \hat{\Lambda}_{\mathcal{V}}} [\mathcal{E}(\hat{\pi})] \right\} &\leq \mathcal{E}^* \\ &+ \min_{\Lambda \in \mathcal{V}} \mathbb{E}_{\pi \sim \Lambda} \left\{ \mathbb{E}_{P_{T+1} \sim \mathcal{P}} \left[\left(\frac{d(P_{T+1}, \pi) \log(n)}{n} \right)^{\beta} \right] \right. \\ &\quad \left. + \frac{\mathcal{K}(\Lambda, \Pi)}{\gamma T} \right\}. \end{aligned}$$

For example, in the case $\mathcal{Q} = \{\pi_1, \dots, \pi_K\}$, taking \mathcal{V} as the set of Dirac masses allows to define $\hat{\pi}$ by a minimization rather than by randomisation.

Note however that our result require to use “exact” Bayes within tasks.

Lemma – Bernstein condition at the meta-level

For any $\pi \in \mathcal{Q}$,

$$\begin{aligned} & \mathbb{E}_{\text{new}} \left[\left(\hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi, \alpha), \pi \right) - \hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi^*, \alpha), \pi^* \right) \right)^2 \right] \\ & \leq C \mathbb{E}_{\text{new}} \left[\hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi, \alpha), \pi \right) - \hat{\mathcal{R}}_{T+1} \left(\rho_{T+1}(\pi^*, \alpha), \pi^* \right) \right]. \end{aligned}$$

We don't know how to extend this lemma if we replace $\rho_{T+1}(\pi, \alpha)$ by a variational approximation.

Some important open questions :

- extending the Lemma to allow variational approximations.
- lower bounds.